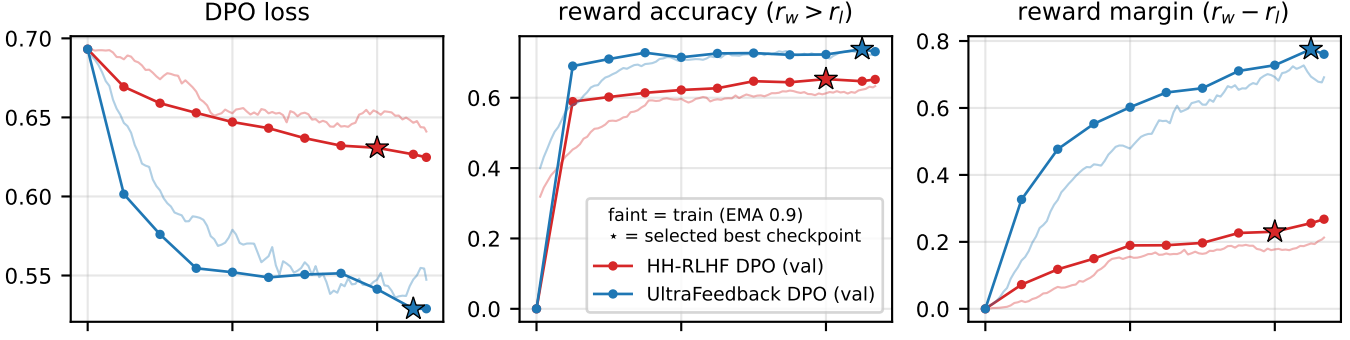
DPO loss
0.70
0.65
0.60
0.55
0
200
400
optimizer step (32 pairs/step)
reward accuracy ($r_w > r_l$)
0.6
0.4
0.2
0.0
faint = train (EMA 0.9)
⋆ = selected best checkpoint
HH-RLHF DPO (val)
UltraFeedback DPO (val)
0
200
400
optimizer step (32 pairs/step)
reward margin ($r_w - r_l$)
0.8
0.6
0.4
0.2
0.0
0
200
400
optimizer step (32 pairs/step)