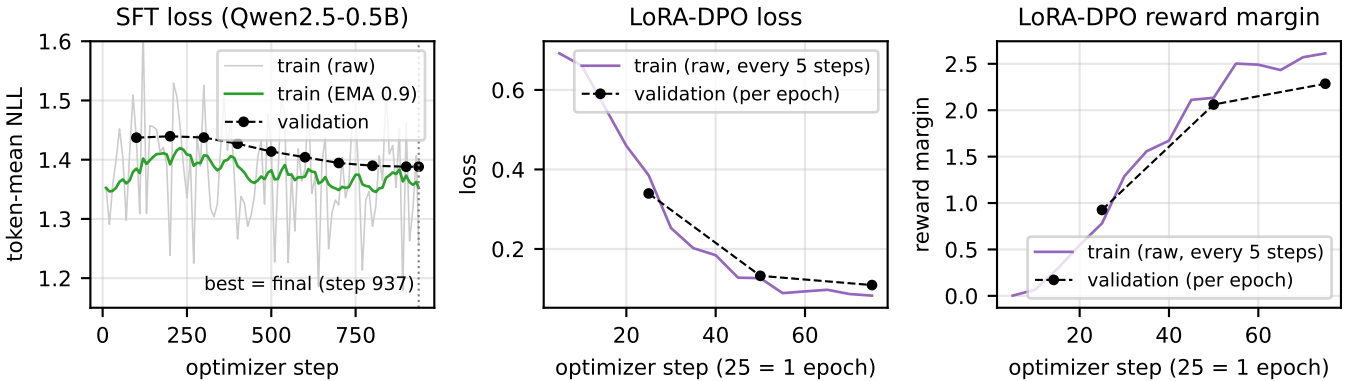
SFT loss (Qwen2.5-0.5B)
token-mean NLL
optimizer step
train (raw)
train (EMA 0.9)
validation
best = final (step 937)
LoRA-DPO loss
loss
optimizer step (25 = 1 epoch)
train (raw, every 5 steps)
validation (per epoch)
LoRA-DPO reward margin
reward margin
optimizer step (25 = 1 epoch)
train (raw, every 5 steps)
validation (per epoch)